MS-II. Concepts

Ali Hashaam, Ali Memon, Guzel Mussilova, Pavlo Shevchenko

Scientific Project: Databases for Multi-Dimensional Data, Genomics and Modern Hardware

May 23, 2017

Table of Contents

Octopus + BlinkDB = Blinktopus

Create a new type of database system without
fixed store that will mimic several existing systems

The goal is to provide approximate answers
with acceptable accuracy in orders of magnitude
less time than that for the exact query processing.[1]

---

[1] Liu, Qing. Approximate Query Processing (Reference work entry) in: Liu, Ling, and M. Tamer zsu.
Encyclopedia of database systems. Vol. 6. Berlin, Heidelberg, Germany: Springer, 2009.

Our Goal

To provide a **framework** that gives user a chance to act as *Holistic SV Optimizer* like in OctopusDB

Add **Approximate Query Processing (AQP)** techniques

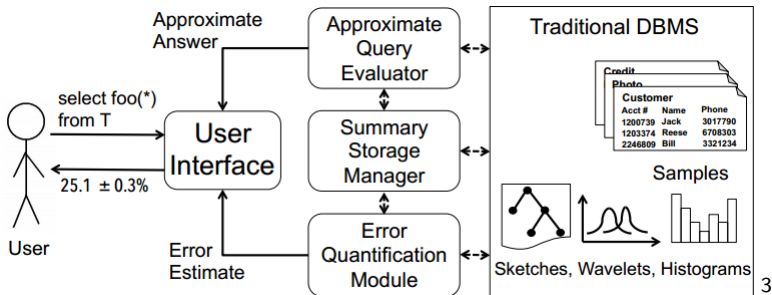**Evaluate** performance depending on choice of SV

Octopus in Blinktopus

(d) SV Transformation: use more efficient Row SV and Col SV [2]

---

[2] Jindal, Alekh. "OctopusDB: flexible and scalable storage management for arbitrary database engines." (2012).

AQP.Architecture

---

[3] The general anatomy of approximate query processing system

AQP.Synopses Manager

A synopsis captures essential properties of the real data while taking less space. The synopses manager is responsible for:

- Type of summary to use(Samples, histograms, sketches, wavelets etc.)
- When to build it (offline vs. online)
- How to store it (to use overlapping samples, how to structure/index/cache the synopses)
- When to update it (batch or online)

Types of Synopses

4 main families of synopses[4]:

- Samples
- Histograms
- Wavelets
- Sketches

---

[4]Cormode, Graham, Minos Garofalakis, Peter J. Haas, and Chris Jermaine. "Synopses for massive data: Samples, histograms, wavelets, sketches." Foundations and Trends in Databases 4, no. 13 (2012): 1-294

Types of Synopses: Samples

Representative subset, chosen by stochastic sampling methods (e.g. Bernoulli, stratified,simple random with and without replacement).

Types of Synopses: Samples

Representative subset, chosen by stochastic sampling methods (e.g. Bernoulli, stratified,simple random with and without replacement).

- Easy to implement.
- Through adhering the same data structure as original tables support the widest range of queries.
- Unbiased estimators for SUM/AVG queries are straightforwardly built.
- Due to the immediate construction after issuing user query do not incur a delay.
- Collecting more samples incrementally enhance imprecise estimates of a query results.
- Without using the attribute value independence assumption allow high quality selectivity estimations.

Types of Synopses: Samples

- Poor estimations for less results.
- Larger relations need more advanced techniques to make the sampling more scalable.
- Selectivity estimations over larger datasets are less efficient.
- Sensitive to skew and outliers.
- Hard to use with (NOT-)/IN, DISTINCT, EXISTS queries.
- Might be difficult to interpret by statistically unsophisticated users.

Types of Synopses: Histograms

A binned representation of the data distribution. The summary and bucket information is used to (approximately) reconstruct the data in the bucket in order to approximately answer the query.

Types of Synopses: Histograms

A binned representation of the data distribution. The summary
and bucket information is used to (approximately) reconstruct the
data in the bucket in order to approximately answer the query.

- A natural solution for range-sum queries.
- Conceptual simplicity allows an effective use of a broad variety
  of estimation tasks(E.g. set-valued queries, real-valued data,
  and aggregate queries over predicates that more complex than
  simple ranges).
- Relatively simple in interpretation.
- Practically acceptable accuracies, provided that the sufficient
  storage space are allocated.

Types of Synopses: Histograms

- Sensitive to dimensionality.
- Performance strongly depends on bucketing schemes(how the buckets are chosen, what statistics are stored, how estimates are extracted, and what classes of query are supported).
- Incremental maintenance.
- Might provide too loose error estimates over the class of queries.

Types of Synopses: Wavelets

Transform the data to represent significant features in a frequency domain and can capture combinations of high and low frequency information.

FAKULTÄT FÜR
INFORMATIK

OTTO VON GUERICKE
UNIVERSITÄT
MAGDEBURG

INF

Types of Synopses: Wavelets

Transform the data to represent significant features in a frequency domain and can capture combinations of high and low frequency information.

- Useful for range-sum queries.

- Appropriately defined AQP algebra allows applying general SPJ (select, project, join) queries on relation summaries.

- The linearity of the basic Haar transform manages a better maintenance under dynamic data than histograms.

- Non-trivial to maintain wavelet coefficients under arbitrary update patterns.

- Large number of coefficients must be retained to guarantee an accurate reconstruction of the data distribution in the multi-dimensional wavelet.

Types of Synopses: Sketches

- Especially appropriate for streaming data.
- Process fasters and easily paralleliazed if each new piece of data is independent of the current state of the summary.
- Can be used as primitives within more complex mining operations, and to extract wavelet and histogram representations of streaming data.

FAKULTÄT FÜR
INFORMATIK

OTTO VON GUERICKE
UNIVERSITÄT
MAGDEBURG

INF

Types of Synopses: Sketches

- Mostly focused on answering a single type of query.
- Number of parameters affects the accuracy and probability of failure.
- Techniques do not extend well to more complex queries which combine multiple sub-queries.
- The only complexity is mathematical (for complete accuracy estimation).

Building a Blinktopus. Recall

First, the Octopus:

- Store incoming data in logs.
- Query the logs (just a filter query)
- Allow users to create views (row, column) over certain logs.
- List all views and logs
- Launch the query over views or over logs, see the changes in performance.

Building a Blinktopus. Recall

Enter AQP:

- What synopsis can we easily support as a view for a specific query? Which will we choose to test? (Samples, histograms?)
- Do Octopuses and AQP match well together?
- How will we allow users to build this view?
- How will we support queries using this view?

Building a Blinktopus. IDE

- Back end

- Front end [5]

---

[5] Sources: http://jupyter.org/
http://honstain.com/new-dropwizard-1-0-5-java-service/

Project Organisation.Roles

Team:

Guzel - Team Leader-Researcher

Pavlo - Developer

Ali H. - Developer

Ali M. - Researcher

Supervisor:

Gabriel Campero Durand

Changing roles after each milestone.

Project Organisation.Schedule

Milestones

Date	Phase	Status
02.05.2017	MS-I (Kick-Off)	done
23.05.2017	MS-II (Concepts)	done
13.06.2017	MS-III (Implementation)	
04.07.2017	MS-IV (Final)	

Meetings

Team Meetings: Mo 14-15
Meetings with supervisor: We 10-11

Thank you! Any questions?

FAKULTÄT FÜR
INFORMATIK

OTTO VON GUERICKE
UNIVERSITÄT
MAGDEBURG

INF

Literature

1. Jindal, Alekh. "The mimicking octopus: Towards a one-size-fits-all database architecture." VLDB PhD Workshop. 2010.

2. Dittrich, Jens, and Alekh Jindal. "Towards a One Size Fits All Database Architecture." CIDR. 2011.

3. Jindal, Alekh. "OctopusDB: flexible and scalable storage management for arbitrary database engines." (2012).

4. Idreos, Stratos, Martin L. Kersten, and Stefan Manegold. "Database Cracking." In CIDR, vol. 7, pp. 68-78. 2007.

5. Mozafari, Barzan. "Approximate query engines: Commercial challenges and research opportunities." SIGMOD, 2017.

6. Cormode, Graham, Minos Garofalakis, Peter J. Haas, and Chris Jermaine. "Synopses for massive data: Samples, histograms, wavelets, sketches." Foundations and Trends in Databases 4, no. 13 (2012): 1-294.