



OTTO VON GUERICKE
UNIVERSITÄT
MAGDEBURG

INF

FAKULTÄT FÜR
INFORMATIK

MS-III. Implementation

Ali Hashaam, Ali Memon, Guzel Mussilova, Pavlo Shevchenko

Scientific Project: Databases for Multi-Dimensional Data, Genomics and Modern Hardware

June 13, 2017

Table of Contents

Blinktopus

Recall

Implementation

- Schema

- OctopusDB

- Approximate Query Processing

 - Histograms

 - Sketches

Project Organisation

- Roles

Literature

Our Goal

To provide a **framework** that gives user a chance to act as *Holistic SV Optimizer* like in OctopusDB

Add **Approximate Query Processing (AQP)** techniques

Evaluate performance depending on choice of SV

Building a Blinktopus. Recall

First, the Octopus:

- Store incoming data in logs.
- Query the logs (just a filter query).
- Allow users to create views (row, column) over certain logs.
- List all views and logs.
- Launch the query over views or over logs, see the changes in performance.

Building a Blinktopus. Recall

Enters Approximate Query Processing (AQP):

- Which synopsis will we choose to test? (Samples, histograms, sketches?)
- Do Octopuses and AQP match well together?
- Build the selected synopsis on the whole data, after data insertions.
- Using the synopsis, answer the user queries by reconstructing the approximate data.

Building a Blinktopus. Implementation

Building a Blinktopus. IDE



Dropwizard

- Back end



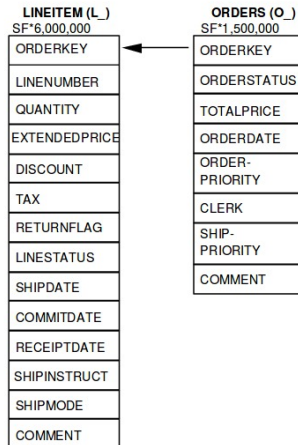
- Front end

1

¹Sources: <http://jupyter.org/>
<http://honstain.com/new-dropwizard-1-0-5-java-service/>

Schema

Selectivity Factors = 1,5,10,15
SF 1 = 1.2 M Records



OctopusDB. Customization




Figure 2: OctopusDB Architecture.

OctopusDB. Customization




Figure 2: OctopusDB Architecture.




Figure 3: Blinktopus.

2

2

A. Jindal. The Mimicking Octopus: Towards a one-size-fits-all Database Architecture, 2010

OctopusDB. Evaluation

Log	190138277,2	8,279069555
RowSV	7001209,38	6,845173066
ColSV	4613638,51	6,664043563




Figure 4: Storage View Type.




Figure 5: Storage View Type (Log Scale).

Evaluation result for 100 runs over Totalprice Column in Orders with Range from 50,000 to 200,000.

AQP. Synopses

4 main families of synopses³:

- Samples
- Histograms ✓
- Wavelets
- Sketches ✓

³Cormode, Graham, Minos Garofalakis, Peter J. Haas, and Chris Jermaine. "Synopses for massive data: Samples, histograms, wavelets, sketches." Foundations and Trends in Databases 4, no. 13 (2012): 1-294. ▶

AQP. Histograms

In histogram's development, main cornerstones are:

- Partition the dataset into buckets.
Number of buckets 'k': $k = 2n^{1/3}$ (RICE RULE)
- Store summary statistics for each bucket about the data values in the it.
- Store information about the buckets themselves, like bucket boundaries.

At query time, the summary and bucket information is used to approximately answer the query.

AQP. Histograms

Vital Points to consider:

- Bucketing Scheme
- Statistics Stored per Bucket
- Approximation Scheme
- Class of queries answered
- Efficiency
- Accuracy & Error Estimates
- Incremental Maintenance

AQP. Histograms

$$D = \{1.61, 1.72, 2.23, 2.33, 2.71, 2.90, 3.41, 4.21, 4.70, 4.82, 4.85, 4.91\}$$





Figure 6: Equi-Depth Histogram ⁴

What if the count of the values between 1.1 and 4.5 is required?

⁴ Cormode, Graham, Minos Garofalakis, Peter J. Haas, and Chris Jermaine. "Synopsis for massive data: Samples, histograms, wavelets, sketches."

AQP. Histograms

Continuous value assumption allows the estimation of values inside a bucket via interpolation.



Actual query : $N = 8$

$AQP : N = 3 + 3 + \left(\frac{4.5-3.4}{4.8-3.4}\right)3 = 8.4$

Overestimation Error = 5%

AQP. Sketches

- Sketches, approximately answer queries by creating small summary data structures that approximately resemble the original data.
- Appropriate in scenarios involving streaming of big data or analysis of higher dimensional data is required.
- Other synopsis (Samples, histograms, Wavelets) can be extracted from it.

AQP. Sketches

Phases of Sketching mechanism:



Figure 7: Sketching Phases⁵

Case Under Scrutiny: Count-Distinct

⁵

Source: <https://yahooeng.tumblr.com/post/135390948446/data-sketches>

AQP. HyperLogLog (HLL)

HLL algorithm estimates the number of distinct elements in large datasets i.e. cardinality, in a single pass, and using a very small amount of memory.⁶

4 billion distinct elements = $\log_2 \log_2(2^{32}) = 5$ bits required

⁶

Flajolet, Philippe, et al. "Understanding the HyperLogLog: a Near-Optimal Cardinality Estimation Algorithm." » « ≡ ≡ ≡ ≡ ≡ ≡ ≡ ≡ ≡

AQP. HyperLogLog (HLL)

INPUT: A Multiset \mathcal{M} of elements from certain domain \mathcal{D} .

OUTPUT: Cardinality estimate n of \mathcal{M} .

1. Define an integer collection C of length $m=2^b$ and initialize its registers with $-\infty$.
2. **for** each element v in \mathcal{M} **do**
3. With a „special“ hash function, translate v to a **stream of bits**, save it in X .
4. With the **first b bits of X** , calculate the **index to update** in collection C . Save it in j .
5. With the **rest of bits of X** , count the **number of 0-bits + 1** at the beginning. Save it in y .
6. Update $C[j]$ only if its actual value is smaller than y . (i.e. $C[j] = \max(C[j], y)$)
7. The current cardinality estimate of n is given by: Harmonic mean of the elements of $C * m * \text{bias correction}$





Figure 8: The HyperLogLog algorithm ⁷

AQP. HyperLogLog (HLL)

- Divide the n distinct elements of the input multiset into m number of buckets.
- Each bucket must comprises approximately the same number of elements, $\frac{n}{m}$.
- transform our input multiset M into an ideal multiset via a special Hash function.
- Hash function will transform the values to a stream of 0s and 1s,

AQP. HyperLogLog (HLL)

- Take first b bits of the hashed value as the index of the bucket.
- In each bucket only save the longest run of starting 0-bits+1 among all the hashed values of the elements that belong to that bucket.

$n = \text{length of longest run of starting 0-bits} + 1$

Number of distinct elements in bucket $= 2^n$

bits required $= \log_2 \log_2(2^n)$

AQP. HyperLogLog (HLL)

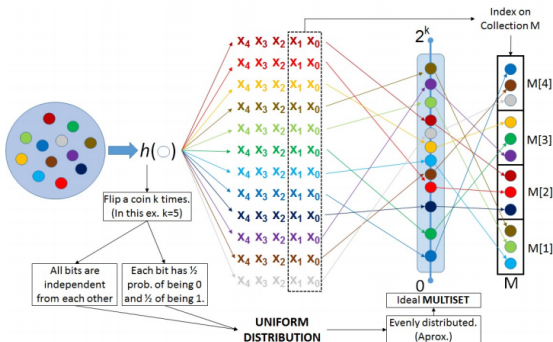


Figure 9: Randomization of elements and division in buckets ⁸

Project Organisation.Roles

Team:

- Guzel - Team Leader-Researcher
- Pavlo - Developer (Backend - OctopusDB)
- Ali H. - Developer (Backend - AQP)
- Ali M. - Developer (Frontend - User Views)

Supervisor:

Gabriel Campero Durand

Changing roles after each milestone.

Meetings:

- Team Meetings: Mo 14-15
- Meetings with supervisor: We 10-11

Thank you! Any questions?

Literature

1. Jindal, Alekh. "The mimicking octopus: Towards a one-size-fits-all database architecture." VLDB PhD Workshop. 2010.
2. Dittrich, Jens, and Alekh Jindal. "Towards a One Size Fits All Database Architecture." CIDR. 2011.
3. Jindal, Alekh. "OctopusDB: flexible and scalable storage management for arbitrary database engines." (2012).
4. Mozafari, Barzan, and Ning Niu. "A Handbook for Building an Approximate Query Engine." IEEE Data Eng. Bull. 38, no. 3 (2015): 3-29.
5. Cormode, Graham, Minos Garofalakis, Peter J. Haas, and Chris Jermaine. "Synopses for massive data: Samples, histograms, wavelets, sketches." Foundations and Trends in Databases 4, no. 13 (2012): 1-294.
6. Flajolet, Philippe, et al. "Understanding the HyperLogLog: a Near-Optimal Cardinality Estimation Algorithm."
7. Flajolet, Philippe, et al. "Hyperloglog: the analysis of a near-optimal cardinality estimation algorithm." Analysis of Algorithms 2007 (AofA07). 2007.